

Homework 1

Jichao Yang

Problem 1

Unless specifically mentioned, the sum and product operators below all range within set $\{x_1, \dots, x_8\}$.

(a) Likelihood function can be found as follow:

$$L = \prod P(X = x_i) = \prod e^{-\lambda} \frac{\lambda^{x_i}}{x_i!} = e^{-8\lambda} \cdot \frac{\lambda^{\sum x_i}}{\prod x_i!}$$

(b) The log likelihood can be found as follow:

$$\mathcal{L} = \log(L) = -8\lambda + \log(\lambda) \cdot \sum x_i - \sum \log(x_i!)$$

(c) Function $\mathcal{L}(\lambda) = c_0 + c_1 \cdot \lambda + c_2 \cdot \ln(\lambda)$ is the sum of convex functions, hence \mathcal{L} is convex. Therefore \mathcal{L} reaches its maximum if and only if its derivative is 0. Solving this equation we have:

$$\begin{aligned} \frac{d}{d\lambda} \mathcal{L} &= -8 + \frac{1}{\lambda} \cdot \sum x_i = 0 \\ \lambda &= \frac{1}{8} \sum x_i \end{aligned}$$

Problem 2

(a) The estimation for λ is

```
import numpy as np

counts = np.array([32,25,28,22,31,34,23,17])
l = sum(counts) / 8

print('The estimation for lambda is {}'.format(np.round(l,5)))
```

The estimation for lambda is 26.5

(b) Below is the function that calculates the negative log likelihood:

```
from scipy.special import factorial

def nll(l, counts):
    '''
    Returns the negative log likelihood of given observation counts and rate
    parameter l.
    '''
    return 8*l - np.log(l)*np.sum(counts) + np.sum(np.log(factorial(counts)))
```

(c) We use the Scipy optimizer to find the minimum negative loss:

```
from scipy.optimize import minimize

# Using initial guess 20 which is decently close to l
l_hat = minimize(nll, [20], args=counts).x[0]
# Does the analytic solution agree with the optimization?
assert np.isclose(l,l_hat)

print('The estimator found using scipy is {}'.format(np.round(l_hat, 5)))
print('The estimator calculated in Problem 1 is {}'.format(np.round(l, 5)))
```

The estimator found using scipy is 26.49999

The estimator calculated in Problem 1 is 26.5

The calculated estimator is reasonably close to the optimized estimator.

Problem 3

(a) Denote the value of the two dice d_1, d_2 . The conditional probability can be calculated with:

$$P(d_1 = 4 | d_1 + d_2 = 7) = \frac{P(d_1 = 4 \cap d_1 + d_2 = 7)}{P(d_1 + d_2 = 7)} = \frac{(1/6)^2}{(1/6)^2 \cdot 6} = \frac{1}{6}$$

(b) Denote families owning dogs D and cats C . The conditional probability can be calculated with:

$$P(D|C) = \frac{P(D \cap C)}{P(C)} = \frac{1/3 \cdot 0.6}{0.4} = \frac{1}{2}$$

Problem 4

- (a) Denote the ordered set of throws $D = \{5, 3, 9, 3, 8, 4, 7\}$. We first calculate a few useful values for solving the problem:

$$P(D|A) = \frac{3 \cdot 3 \cdot 2}{20^7} = \frac{18}{20^7}; \quad P(D|B) = \frac{2^6}{20^7} = \frac{64}{20^7}; \quad P(D|E) = \frac{2^7}{20^7} = \frac{128}{20^7}$$

$$P(D) = P(D|A) \cdot P(A) + P(D|B) \cdot P(B) + P(D|E) \cdot P(E) = \frac{1}{3} \cdot \frac{210}{20^7}$$

Using the Bayes equation we have the following:

$$P(A|D) = \frac{P(D|A) \cdot P(A)}{P(D)} = \frac{18}{210}; \quad P(B|D) = \frac{64}{210}; \quad P(E|D) = \frac{128}{210}$$

- (b) Take calculations done in step (a) as the prior. Denote D' as the final die roll of 10. We can calculate the following:

$$P(D') = P(D'|A)P(A) + P(D'|B)P(B) + P(D'|E)P(E) = 0 \cdot \frac{18}{210} + \frac{1}{20} \cdot \frac{64}{210} + \frac{2}{20} \cdot \frac{128}{210} = \frac{16}{210}$$

$$P(A|D') = \frac{P(D'|A) \cdot P(A)}{P(D')} = \frac{0 \cdot 18/210}{16/210} = 0; \quad P(B|D') = \frac{1}{5}; \quad P(E|D') = \frac{4}{5}$$

- (c) The consequences of A being impossible to roll a 10 is that the posterior probability of A is 0. In other words, the prior probability of A becomes ‘distributed’ among B and C after the additional roll.
 (d) Recall that given a set of rolls D , there is a established 99:1 confidence between B and E if and only if

$$\frac{P(B|D)}{P(E|D)} \leq \frac{1}{99} \text{ or } \frac{P(B|D)}{P(E|D)} \geq \frac{99}{1}$$

Given any prior rolls D , we can calculate the current value for $P(B)$ and $P(E)$. For new roll D' , the probability can be updated as follow:

$$P(B|D') = \frac{P(D'|B)}{P(D')} \cdot P(B); \quad P(E|D') = \frac{P(D'|E)}{P(D')} \cdot P(E)$$

Hence the quotient can be updated as follow:

$$\frac{P(B|D')}{P(E|D')} = \frac{P(D'|B)}{P(D'|E)} \cdot \frac{P(B)}{P(E)}$$

For accuracy of calculation, we define the log confidence as

$$\mathcal{L} = \log\left(\frac{P(B)}{P(E)}\right)$$

Hence the log confidence can be iteratively updated by the following rule

$$\mathcal{L}' = \mathcal{L} + \log(P(D'|B)) - \log(P(D'|E))$$

until we reach the breaking condition

$$\mathcal{L}' \leq -\log(99) \text{ or } \mathcal{L}' \geq \log(99)$$

Below is the implementation of the algorithm:

```

import random
from tqdm import tqdm
import matplotlib.pyplot as plt

# Set up dice rolls
B = [1,1,1,2,2,2,3,3,4,4,5,5,6,6,7,7,8,8,9,10]
E = [1,1,2,2,3,3,4,4,5,5,6,6,7,7,8,8,9,9,10,10]
B_chance = {
    1: 3/20, 2: 3/20, 3: 2/20, 4: 2/20, 5: 2/20,
    6: 2/20, 7: 2/20, 8: 2/20, 9: 1/20, 10: 1/20
}
E_chance = {
    1: 2/20, 2: 2/20, 3: 2/20, 4: 2/20, 5: 2/20,
    6: 2/20, 7: 2/20, 8: 2/20, 9: 2/20, 10: 2/20
}

def simulate():
    """
    Returns the needed number of rolls to establish 99:1
    confidence for one simulation.
    """
    # Set up simulation parameters
    loss = 0 # assume equal starting chance, hence log(1)
    roll_count = 0
    is_B = random.choice([True, False]) # Pick a die at random

    # Begin simulation
    while(abs(loss) <= np.log(99)):
        # Roll the die
        roll_val = random.choice(B) if is_B else random.choice(E)
        # Update loss
        loss = loss + np.log(B_chance[roll_val]) - np.log(E_chance[roll_val])
        roll_count += 1
    return roll_count

# Find out the mean of the necessary roll counts
n = 10000; results = []
for i in range(n):
    results.append(simulate())
mean = np.mean(results)
plt.hist(results, bins=50)
plt.axvline(mean, color='k', linestyle='dashed', linewidth=1)
plt.text(mean*1.1, plt.ylim()[1]*0.9, 'Mean: {:.0f}'.format(mean))
plt.xlabel('Number of necessary rolls for 99:1 confidence')
plt.ylabel('Count')
print('It takes {:.0f} rolls on average to establish a 99:1 confidence'.format(mean))
plt.show()

```

It takes 88 rolls on average to establish a 99:1 confidence




Figure 1: Number of rolls necessary on average for 99:1 confidence

Problem 5

- (a) The probability of making k of 18 throws is

$$P(k|p) = \frac{18!}{(18-k)! \cdot k!} \cdot p^k \cdot (1-p)^{18-k}$$

- (b) The likelihood function for π is

$$\begin{aligned} P(\pi|k) &= \frac{P(k|\pi) \cdot P(\pi)}{P(k)} \sim P(k|\pi) \\ &\sim p^k \cdot (1-p)^{18-k} \end{aligned}$$

- (c) Observe that the likelihood function for π is convex in $[0, 1]$. The negative log likelihood function \mathcal{L} as

$$\mathcal{L} = -\log(p^k \cdot (1-p)^{18-k}) = (-k) \cdot \log(p) + (k-18) \cdot \log(1-p)$$

Given convexity, to find the minimum of \mathcal{L} , we find the zero for its derivative.

$$\begin{aligned} \frac{d\mathcal{L}}{dp} &= -\frac{k}{p} + \frac{18-k}{1-p} = 0 \\ \hat{p} &= \frac{k}{18} \end{aligned}$$

- (d) Given $k = 12$, we can calculate

$$\hat{p} = \frac{12}{18} = \frac{2}{3}$$